

# Marginal means with lmer

Sasha D. Hafner

18 November, 2024 Nov:11

## Summary

Just confirming that emmeans does what I think it does with mixed-effects models without balance.

## 0. Packages

```
library(lme4)
library(emmeans)
library(data.table)
```

```
source('dfcombos.R')
```

## 1. Data

Five levels of a random effect.

```
rv <- data.table(rv = factor(letters[1:5]), re = -2:2)
rv
```

```
##           rv      re
##    <fctr> <int>
```

```
## 1:      a      -2
## 2:      b      -1
## 3:      c       0
## 4:      d       1
## 5:      e       2
```

(Hmm, should probably try with a mean different from 0.)

Three treatments make up a fixed effect.

```
ff <- data.table(trt = factor(letters[1:3]), fe = c(-1, 0, 1))
```

Combine.

```
d0 <- dfcombos(rv, ff)
```

Eliminate balance.

```
d0 <- d0[-c(9:10, 12:14), ]
d0
```

```
##      rv re trt fe
## 1     a -2   a -1
## 2     b -1   a -1
## 3     c  0   a -1
## 4     d  1   a -1
## 5     e  2   a -1
## 6     a -2   b  0
## 7     b -1   b  0
## 8     c  0   b  0
## 11    a -2   c  1
## 15    e  2   c  1
```

```
table(d0$trt)
```

```
##
## a b c
## 5 3 2
```

```
table(d0$rv)
```

```
##  
## a b c d e  
## 3 2 2 1 2
```

Add response variable

```
setDT(d0)  
d0[, y := re + fe + rnorm(nrow(d0), sd = 0.2)]  
d0
```

```
##      rv    re   trt   fe      y  
##   <fctr> <int> <fctr> <num>   <num>  
## 1:     a    -2     a    -1 -2.85217702  
## 2:     b    -1     a    -1 -1.92267825  
## 3:     c     0     a    -1 -0.74072057  
## 4:     d     1     a    -1 -0.16071167  
## 5:     e     2     a    -1  0.67947487  
## 6:     a    -2     b     0 -1.81334981  
## 7:     b    -1     b     0 -0.63878215  
## 8:     c     0     b     0 -0.01130073  
## 9:     a    -2     c     1 -0.62281774  
## 10:    e     2     c     1  3.31567669
```

## 2. Model

```
m1 <- lmer(y ~ trt + (1|rv), data = d0)  
summary(m1)
```

```
## Linear mixed model fit by REML ['lmerMod']  
## Formula: y ~ trt + (1 | rv)  
## Data: d0  
##
```

```
## REML criterion at convergence: 18.6
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -0.7665 -0.5892  0.1489  0.4727  0.7765
##
## Random effects:
##   Groups   Name      Variance Std.Dev.
##   rv       (Intercept) 2.04697  1.4307
##   Residual              0.04147  0.2036
## Number of obs: 10, groups: rv, 5
##
## Fixed effects:
##              Estimate Std. Error t value
## (Intercept)  -0.9994     0.6463  -1.546
## trtb         1.0436     0.1622   6.433
## trtc         2.4095     0.1938  12.431
##
## Correlation of Fixed Effects:
##      (Intr) trtb
## trtb -0.079
## trtc -0.066  0.240
```

```
emmeans(m1, 'trt')
```

```
##   trt   emmean    SE    df lower.CL upper.CL
## a  -0.9994 0.646 4.06   -2.784    0.785
## b   0.0443 0.654 4.24   -1.731    1.820
## c   1.4101 0.663 4.45   -0.358    3.178
##
## Degrees-of-freedom method: kenward-roger
## Confidence level used: 0.95
```

Compare to mean by treatment.

```
d0[, mean(y), by = trt]
```

```
##      trt      V1
```

```
##      <fctr>      <num>
## 1:      a -0.9993625
## 2:      b -0.8211442
## 3:      c  1.3464295
```

### 3. Higher n

Instead of running this many times, I'll just increase the sample size.

```
d1 <- d0[rep(1:nrow(d0), each = 300), ]
table(d1$trt)
```

```
##
##      a      b      c
## 1500  900  600
```

And recreate y values.

```
set.seed(1)
d1[, y := re + fe + rnorm(nrow(d1), sd = 0.2)]
d1
```

```
##      rv      re      trt      fe      y
##      <fctr> <int> <fctr> <num>      <num>
## 1:      a      -2      a      -1 -3.125291
## 2:      a      -2      a      -1 -2.963271
## 3:      a      -2      a      -1 -3.167126
## 4:      a      -2      a      -1 -2.680944
## 5:      a      -2      a      -1 -2.934098
## ---
## 2996:      e      2      c      1  2.962648
## 2997:      e      2      c      1  2.954128
## 2998:      e      2      c      1  3.326037
## 2999:      e      2      c      1  2.567066
## 3000:      e      2      c      1  2.784445
```

```
m2 <- lmer(y ~ trt + (1|rv), data = d1)
summary(m2)
```

```
## Linear mixed model fit by REML ['lmerMod']
## Formula: y ~ trt + (1 | rv)
## Data: d1
##
## REML criterion at convergence: -870.8
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -3.4077 -0.6384 -0.0169  0.6887  3.6976
##
## Random effects:
## Groups Name Variance Std.Dev.
## rv      (Intercept) 2.5117  1.5848
## Residual 0.0429  0.2071
## Number of obs: 3000, groups: rv, 5
##
## Fixed effects:
##              Estimate Std. Error t value
## (Intercept) -1.001752  0.708775  -1.413
## trtb         0.999657  0.009539 104.798
## trtc         2.001700  0.011401 175.571
##
## Correlation of Fixed Effects:
##      (Intr) trtb
## trtb -0.004
## trtc -0.004  0.239
```

```
emmeans(m2, 'trt')
```

```
## trt  emmean    SE df lower.CL upper.CL
## a   -1.0018 0.709  4   -2.970    0.966
## b   -0.0021 0.709  4   -1.970    1.966
## c    0.9999 0.709  4   -0.968    2.968
##
```

```
## Degrees-of-freedom method: kenward-roger  
## Confidence level used: 0.95
```

Compare to mean by treatment.

```
d1[, mean(y), by = trt]
```

```
##      trt      V1  
##   <fctr>   <num>  
## 1:      a -1.001752  
## 2:      b -1.003417  
## 3:      c  1.005306
```

Notice that `lmer` coefficients (`Estimate`) and marginal means look identical.

## Conclusion

The `emmeans` package seems to be a simple way to get marginal means from mixed-effects models.